

data_analysis

June 29, 2025

```
[1]: import pyspark.sql.functions as sf
import pandas as pd
from pyspark.sql import SparkSession
import matplotlib.pyplot as plt
```

```
[2]: # Create pandas DataFrame
df_pd = pd.read_csv("metadata_output/combined_metadata.csv")
```

```
[3]: # Create Spark DataFrame
spark = SparkSession.builder.appName("Data Analysis App").getOrCreate()
df_sp = spark.read.csv('metadata_output/combined_metadata.csv', header=True,
    ↪inferSchema=True)
```

3.1 Descriptive Statistics

```
[4]: # 1. What is the average duration (in seconds) of all videos in the dataset?
average_duration_pd = df_pd["duration_seconds"].mean()
print(f"[Pandas] Average duration of all videos: {average_duration_pd:.2f}
    ↪seconds")

average_duration_sp = df_sp.select(sf.mean("duration_seconds"))
print("[Spark] Average duration of all videos:")
average_duration_sp.show()
```

[Pandas] Average duration of all videos: 1508.36 seconds

[Spark] Average duration of all videos:

```
+-----+
|avg(duration_seconds)|
+-----+
| 1508.3636363636363|
+-----+
```

```
[5]: # 2. Which uploader appears most frequently in the dataset?
mode_uploader_pd = df_pd["uploader"].mode()
print(f"[Pandas] Most frequent uploader: {mode_uploader_pd[0]}")

mode_uploader_sp = df_sp.groupBy("uploader").count().orderBy(sf.desc("count")).
    ↪limit(1)
```

```
print("[Spark] Most frequent uploader:")
mode_uploader_sp.show()
```

[Pandas] Most frequent uploader: Chappell Roan

[Spark] Most frequent uploader:

```
+-----+-----+
|      uploader|count|
+-----+-----+
|Chappell Roan|    2|
+-----+-----+
```

[6]: *# 3. Which five videos have the highest number of views? List their titles and view counts.*

```
top_views_pd = df_pd.nlargest(5, "view_count")[["title", "view_count"]]
print(f"[Pandas] Top 5 videos by view count: {top_views_pd.
      ↪to_string(index=False)}")
```

```
top_views_sp = df_sp.select("title", "view_count").orderBy(sf.
      ↪desc("view_count")).limit(5)
print("[Spark] Top 5 videos by view count:")
top_views_sp.show()
```

[Pandas] Top 5 videos by view count:

title	view_count
Chappell Roan - Pink Pony Club (Official Music Video)	78618729
Fleetwood Mac - Silver Springs (Live) (Official Video) [HD]	48786565
Mozart - Classical Music for Brain Power	41299056
Charlie Chaplin - Final Speech from The Great Dictator	9749956
Chappell Roan - The Giver (Official Lyric Video)	7319955

[Spark] Top 5 videos by view count:

```
+-----+-----+
|      title|view_count|
+-----+-----+
|Chappell Roan - P...| 78618729|
|Fleetwood Mac - S...| 48786565|
|Mozart - Classica...| 41299056|
|Charlie Chaplin -...|  9749956|
|Chappell Roan - T...|  7319955|
+-----+-----+
```

[7]: *# 4. For each upload year, what is the average number of likes?*

```
average_likes_per_year_pd = df_pd.groupby("year_uploaded")["like_count"].mean().
      ↪reset_index()
average_likes_per_year_sp = average_likes_per_year_pd.
      ↪sort_values(by="like_count", ascending=False)
```

```

print(f"[Pandas] Average likes per upload year:\n{average_likes_per_year_pd.
    ↳to_string(index=False)}")

average_likes_per_year_sp = df_sp.groupBy("year_uploaded").agg(sf.
    ↳mean("like_count").alias("average_likes"))
average_likes_per_year_sp = average_likes_per_year_sp.orderBy(sf.
    ↳desc("average_likes"))
print("[Spark] Average likes per upload year:")
average_likes_per_year_sp.show()

```

[Pandas] Average likes per upload year:

year_uploaded	like_count
2020	704145.0
2018	336813.0
2016	270753.0
2025	108272.0
2022	19747.5
2023	3899.5
2019	811.0

[Spark] Average likes per upload year:

year_uploaded	average_likes
2020	704145.0
2018	336813.0
2016	270753.0
2025	108272.0
2022	19747.5
2023	3899.5
2019	811.0

[8]: # 5. How many videos are missing artist information?

```

missing_artist_count_pd = df_pd["artist"].isnull().sum()
print(f"[Pandas] Number of videos missing artist information:␣
    ↳{missing_artist_count_pd}")

missing_artist_count_sp = df_sp.filter(sf.col("artist").isNull()).count()
print(f"[Spark] Number of videos missing artist information:␣
    ↳{missing_artist_count_sp}")

```


[Pandas] Number of videos missing artist information: 11

[Spark] Number of videos missing artist information: 11

3.2 Tag and content characteristics

[9]: # 1. How many tags does each video have? Visualize the distribution using a `histogram`.

```
plt.figure(figsize=(6, 4))
bars = plt.bar(df_pd["legible_title"], df_pd["tag_count"])
plt.title("Tag count distribution")
plt.xlabel("Video")
plt.ylabel("Number of Tags")
plt.xticks(rotation=45, ha='right')
plt.bar_label(
    bars,
    fmt='%d',
    label_type='center',
    color='white',
    fontsize=9,
    fontweight='bold'
)
plt.tight_layout()
plt.show()
```



[10]: # 2. What is the total number of views per uploader? Rank the results in `descending order`.

```
total_views_per_uploader_pd = df_pd.groupby("uploader").
    agg(total_views=("view_count", "sum"))
```

```

).reset_index()
total_views_per_uploader_pd = total_views_per_uploader_pd.
    ↪sort_values(by="total_views", ascending=False)
print(f"[Pandas] Total views per uploader:\n{total_views_per_uploader_pd.
    ↪to_string(index=False)}")

total_views_per_uploader_sp = df_sp.groupBy("uploader").agg(sf.
    ↪sum("view_count").alias("total_views"))
total_views_per_uploader_sp = total_views_per_uploader_sp.orderBy(sf.
    ↪desc("total_views"))
print("[Spark] Total views per uploader:")
total_views_per_uploader_sp.show()

```

```

[Pandas] Total views per uploader:
      uploader  total_views
Chappell Roan    85938684
Fleetwood Mac    48786565
HALIDONMUSIC     41299056
Charlie Chaplin   9749956
STILL I RISE Motivation  2270465
Saturday Night Live  1266810
Yard Act          826491
lizlawrencemusic   49600
FOX 29 Philadelphia  30102
Mega Shortz        5520

```

```

[Spark] Total views per uploader:
+-----+-----+
|      uploader|total_views|
+-----+-----+
|    Chappell Roan|    85938684|
|    Fleetwood Mac|    48786565|
|    HALIDONMUSIC|    41299056|
|    Charlie Chaplin|    9749956|
|STILL I RISE Moti...|    2270465|
|Saturday Night Live|    1266810|
|        Yard Act|    826491|
|    lizlawrencemusic|    49600|
|FOX 29 Philadelphia|    30102|
|        Mega Shortz|    5520|
+-----+-----+

```

```

[11]: # 3. Which video has the longest duration? List the title and its duration.
longest_video_pd = df_pd.loc[df_pd["duration_seconds"].idxmax(), ["title",
    ↪"duration_seconds"]]
print(f"[Pandas] Longest video: {longest_video_pd['title']} with duration_
    ↪{longest_video_pd['duration_seconds']} seconds")

```

```

longest_video_sp = df_sp.orderBy(sf.desc("duration_seconds")).select("title",
↳ "duration_seconds").limit(1)
print("[Spark] Longest video:")
longest_video_sp.show()

```

[Pandas] Longest video: Mozart - Classical Music for Brain Power with duration 7342 seconds

```

[Spark] Longest video:
+-----+-----+
|          title|duration_seconds|
+-----+-----+
|Mozart - Classica...|          7342|
+-----+-----+

```

[12]: # 4. How many videos were uploaded in each year? Present the results sorted by
↳ year.

```

videos_per_year_pd = df_pd["year_uploaded"].value_counts().reset_index()
videos_per_year_pd.columns = ["year_uploaded", "video_count"]
videos_per_year_pd = videos_per_year_pd.sort_values(by="year_uploaded")
print(f"[Pandas] Number of videos uploaded per year:\n{videos_per_year_pd.
↳ to_string(index=False)}")

```

```

videos_per_year_sp = df_sp.groupBy("year_uploaded").count().
↳ orderBy("year_uploaded")
print("[Spark] Number of videos uploaded per year:")
videos_per_year_sp.show()

```

[Pandas] Number of videos uploaded per year:

year_uploaded	video_count
2016	1
2018	2
2019	1
2020	1
2022	2
2023	2
2025	2

[Spark] Number of videos uploaded per year:

```

+-----+-----+
|year_uploaded|count|
+-----+-----+
|          2016|    1|
|          2018|    2|
|          2019|    1|
|          2020|    1|
|          2022|    2|
|          2023|    2|

```

```
[13]: # 5. Is there a correlation between the number of views and the number of likes?
      ↪ Feel free to drop or filter rows with missing or zero values before
      ↪ computing correlation.
correlation = df_pd[["view_count", "like_count"]].dropna().corr().iloc[0, 1]
if correlation > 0.8:
    assessment = "strong positive correlation"
elif correlation < -0.8:
    assessment = "strong negative correlation"
else:
    assessment = "weak or no correlation"
print(f"[Pandas] The correlation between views and likes is {correlation:.2f},
      ↪ this is a {assessment}.")

correlation = df_sp.select("view_count", "like_count").na.drop()
correlation = correlation.stat.corr("view_count", "like_count")
print(f"[Spark] The correlation between views and likes is {correlation:.2f},
      ↪ this is a {assessment}.")
```

[Pandas] The correlation between views and likes is 0.94, this is a strong positive correlation.

[Spark] The correlation between views and likes is 0.94, this is a strong positive correlation.

3.3 Derived Metrics & Custom Analysis

```
[14]: # 1. Which video has the highest number of likes per second of duration?
highest_likes_per_second_pd = df_pd.assign(
    likes_per_second=lambda x: x["like_count"] / x["duration_seconds"]
).loc[df_pd["like_count"] > 0].nlargest(1, "likes_per_second")[["title",
    ↪ "likes_per_second"]]
print(f"[Pandas] Video with highest likes per second:
      ↪ {highest_likes_per_second_pd['title'].values[0]} with
      ↪ {highest_likes_per_second_pd['likes_per_second'].values[0]:.2f} likes/second.
      ↪ ")

highest_likes_per_second_sp = df_sp.withColumn(
    "likes_per_second", sf.col("like_count") / sf.col("duration_seconds"))
highest_likes_per_second_sp = highest_likes_per_second_sp.filter(sf.
    ↪ col("like_count") > 0)
highest_likes_per_second_sp = highest_likes_per_second_sp.orderBy(sf.
    ↪ desc("likes_per_second")).select("title", "likes_per_second").limit(1)
print("[Spark] Video with highest likes per second:")
highest_likes_per_second_sp.show()
```

[Pandas] Video with highest likes per second: Chappell Roan - Pink Pony Club

(Official Music Video) with 2514.80 likes/second.

[Spark] Video with highest likes per second:

```
+-----+-----+
|          title| likes_per_second|
+-----+-----+
|Chappell Roan - P...|2514.8035714285716|
+-----+-----+
```

[15]: *# 2. Which uploader has the longest total duration of all their uploaded videos, combined?*

```
longest_total_duration_uploader_pd = df_pd.
    ↳groupby("uploader")["duration_seconds"].sum().reset_index()
longest_total_duration_uploader_pd = longest_total_duration_uploader_pd.
    ↳loc[longest_total_duration_uploader_pd["duration_seconds"].idxmax()]
print(f"[Pandas] Uploader with longest total duration:
    ↳{longest_total_duration_uploader_pd['uploader']} with
    ↳{longest_total_duration_uploader_pd['duration_seconds']} seconds.")

longest_total_duration_uploader_sp = df_sp.groupBy("uploader").agg(sf.
    ↳sum("duration_seconds").alias("total_duration"))
longest_total_duration_uploader_sp = longest_total_duration_uploader_sp.
    ↳orderBy(sf.desc("total_duration")).select("uploader", "total_duration").
    ↳limit(1)
print("[Spark] Uploader with longest total duration:")
longest_total_duration_uploader_sp.show()
```

[Pandas] Uploader with longest total duration: HALIDONMUSIC with 7342 seconds.

[Spark] Uploader with longest total duration:

```
+-----+-----+
|  uploader|total_duration|
+-----+-----+
|HALIDONMUSIC|          7342|
+-----+-----+
```

[16]: *# 3. What is the ratio of views to likes for each video?*

```
views_to_likes_ratio_pd = df_pd.assign(
    views_to_likes_ratio=lambda x: x["view_count"] / x["like_count"]
).loc[df_pd["like_count"] > 0][["title", "views_to_likes_ratio"]]
print(f"[Pandas] Views to likes ratio for each video:\n{views_to_likes_ratio_pd.
    ↳to_string(index=False)}")

views_to_likes_ratio_sp = df_sp.withColumn(
    "views_to_likes_ratio", sf.col("view_count") / sf.col("like_count"))
views_to_likes_ratio_sp = views_to_likes_ratio_sp.filter(sf.col("like_count") >
    ↳0)
```



```
views_to_likes_ratio_sp = views_to_likes_ratio_sp.select("title",
↳"views_to_likes_ratio")
print("[Spark] Views to likes ratio for each video:")
views_to_likes_ratio_sp.show(truncate=False)
```

[Pandas] Views to likes ratio for each video:

	title
views_to_likes_ratio	
111.651335	Chappell Roan - Pink Pony Club (Official Music Video)
33.836360	Chappell Roan - The Giver (Official Lyric Video)
36.010519	Charlie Chaplin - Final Speech from The Great Dictator
133.349091	Fleetwood Mac - Silver Springs (Live) (Official Video) [HD]
61.159063	Liz Lawrence - None Of My Friends
134.188050	Mozart - Classical Music for Brain Power
81.995847	One of The Greatest Speeches Ever by President Obama Best Eye Opening Speech
262.857143	Oprah Winfrey Interview With Prince Harry and Meghan Markle Full Interview
162.870918	Saturday Night Live Duos - SNL
143.342857	Trump swearing video: Says Iran, Israel 'don't know what the f-' they're doing
70.011944	Yard Act - 100% Endurance

[Spark] Views to likes ratio for each video:

```
+-----+
+-----+
|title
|views_to_likes_ratio|
+-----+
+-----+
|Chappell Roan - Pink Pony Club (Official Music Video)
|111.65133459727755 |
|Chappell Roan - The Giver (Official Lyric Video)
|33.836359518152484 |
|Charlie Chaplin - Final Speech from The Great Dictator
|36.01051881234926 |
|Fleetwood Mac - Silver Springs (Live) (Official Video) [HD]
|133.34909089915158 |
|Liz Lawrence - None Of My Friends
|61.159062885326755 |
|Mozart - Classical Music for Brain Power
```

|134.18804951749684 |
|One of The Greatest Speeches Ever by President Obama | Best Eye Opening
Speech|81.99584687612857 |
|Oprah Winfrey Interview With Prince Harry and Meghan Markle Full Interview
|262.85714285714283 |
|Saturday Night Live Duos - SNL
|162.87091797377218 |
|Trump swearing video: Says Iran, Israel 'don't know what the f-' they're
doing|143.34285714285716 |
|Yard Act - 100% Endurance
|70.01194409148665 |
+-----+
+-----+

[]: